

MTRX5700 - Assignment Task 2

Report

Experimental Robotics

SIDs:
470148329
470416309
480436153

Question 1 - Hand Eye Calibration	3
Theoretical Analyses	3
Work on RVIZ	4
Question 2 - Point Cloud Registration	7
Point Cloud Poses	7
MATLAB Two Point Cloud Visualisation	8
Results	9
Motion of the Camera	9
Point Cloud Registration	10
Computation of Errors	10
Question 3 - Image Classification using Deep Learning	11
Initial Experimenting	11
Deep Nerual Network	11
Confusion Matrix	11
External Input	13
Question 4 - Perception for Robots	14
Process of Navigation	14
Question 5 - Demo	15
Appendix	16
Codes for Question 3-4	16
Codes for Question 5	23

Question 1 - Hand Eye Calibration

Theoretical Analyses

The diagram illustrates a hand-eye calibration setup. A vertical line labeled 'B' represents the base of a robotic arm. A horizontal line labeled 'H' extends from the end of the arm. A camera, labeled 'E - camera', is positioned at the end of the arm, looking towards a calibration pattern, labeled 'T - calibration Pattern'. Below the diagram, the transformation equation is derived:

$${}^B T_T = \frac{{}^B T_H}{\text{constant}} \cdot \frac{{}^H T_E}{\text{variable}} \cdot \frac{{}^E T_T}{\text{constant}}$$

A downward arrow indicates the simplification of the equation:

$${}^B T_{H_1} \cdot \frac{{}^H T_E}{\text{variable}} \cdot {}^E T_{T_1} = {}^B T_{H_2} \cdot \frac{{}^H T_E}{\text{variable}} \cdot {}^E T_{T_2}$$

The equation is further simplified by canceling out the common terms:

$$\left({}^B T_{H_2} + {}^B T_{H_1} \right) \frac{{}^H T_E}{\text{variable}} = \frac{{}^H T_E}{\text{variable}} \left({}^E T_{T_2} + {}^E T_{T_1} \right)$$

A note at the bottom states: "could be solved in RVIZ simulation".

Figure 1 - Theoretical Analysis

Work on RVIZ

Based on the RVIZ, firstly, setting up the parameters of HandEye calibration and then set UR5e arm in the middle position as shown below and taking a sample:

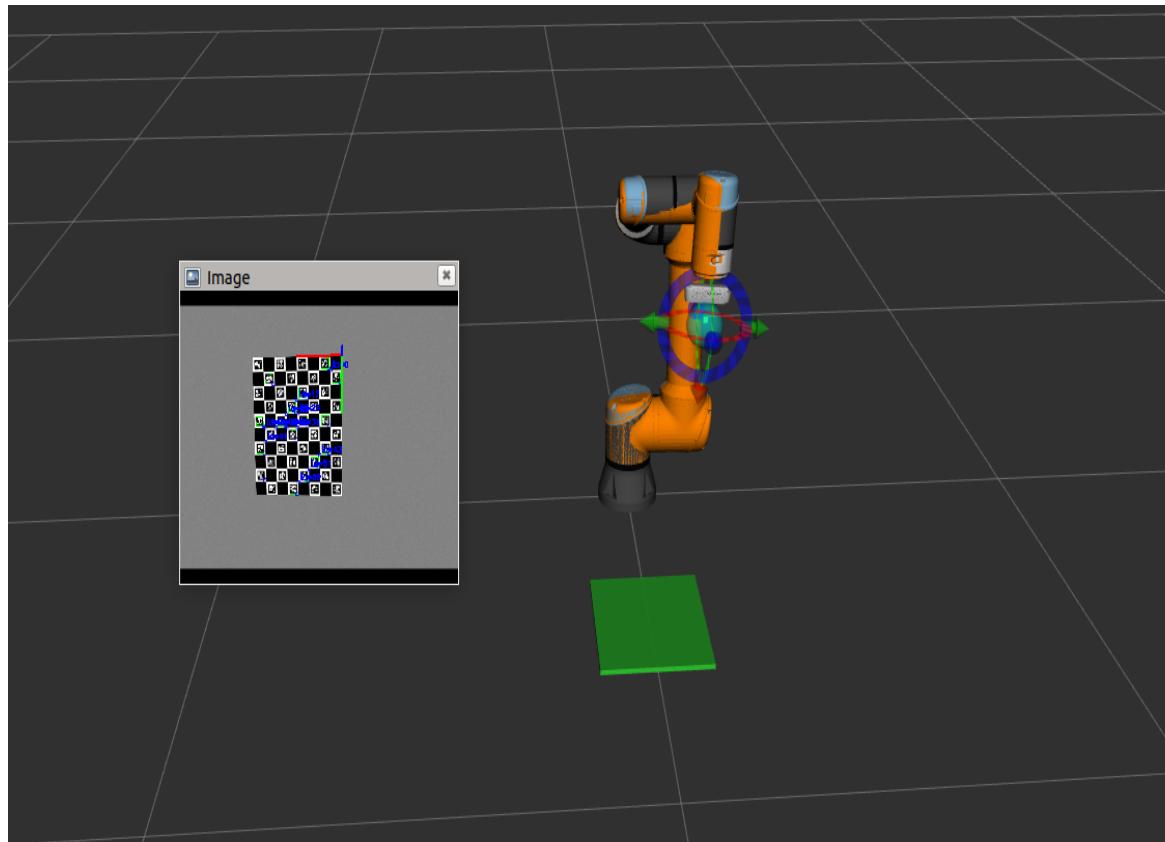
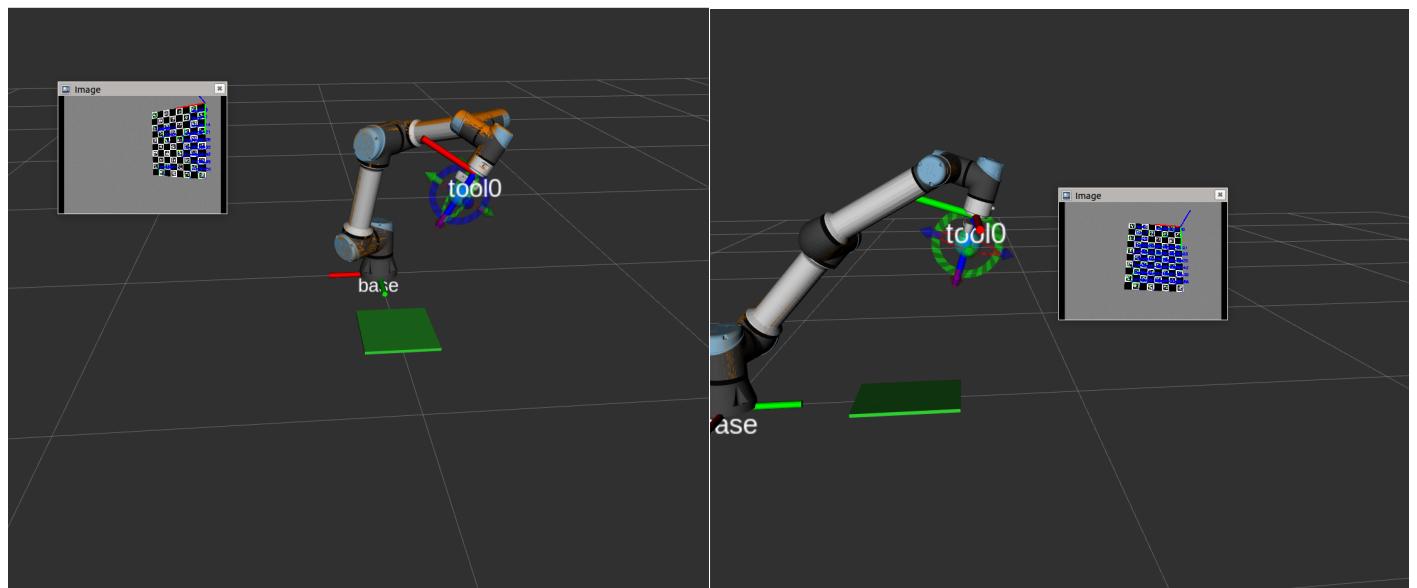


Figure 2 - Located the UR5 Arm in the Middle Position

Second is move the arm in different locations but need to keep the calibration pattern fully observed and some special and key positions are shown below and 30 samples have been taken to increase the accuracy.



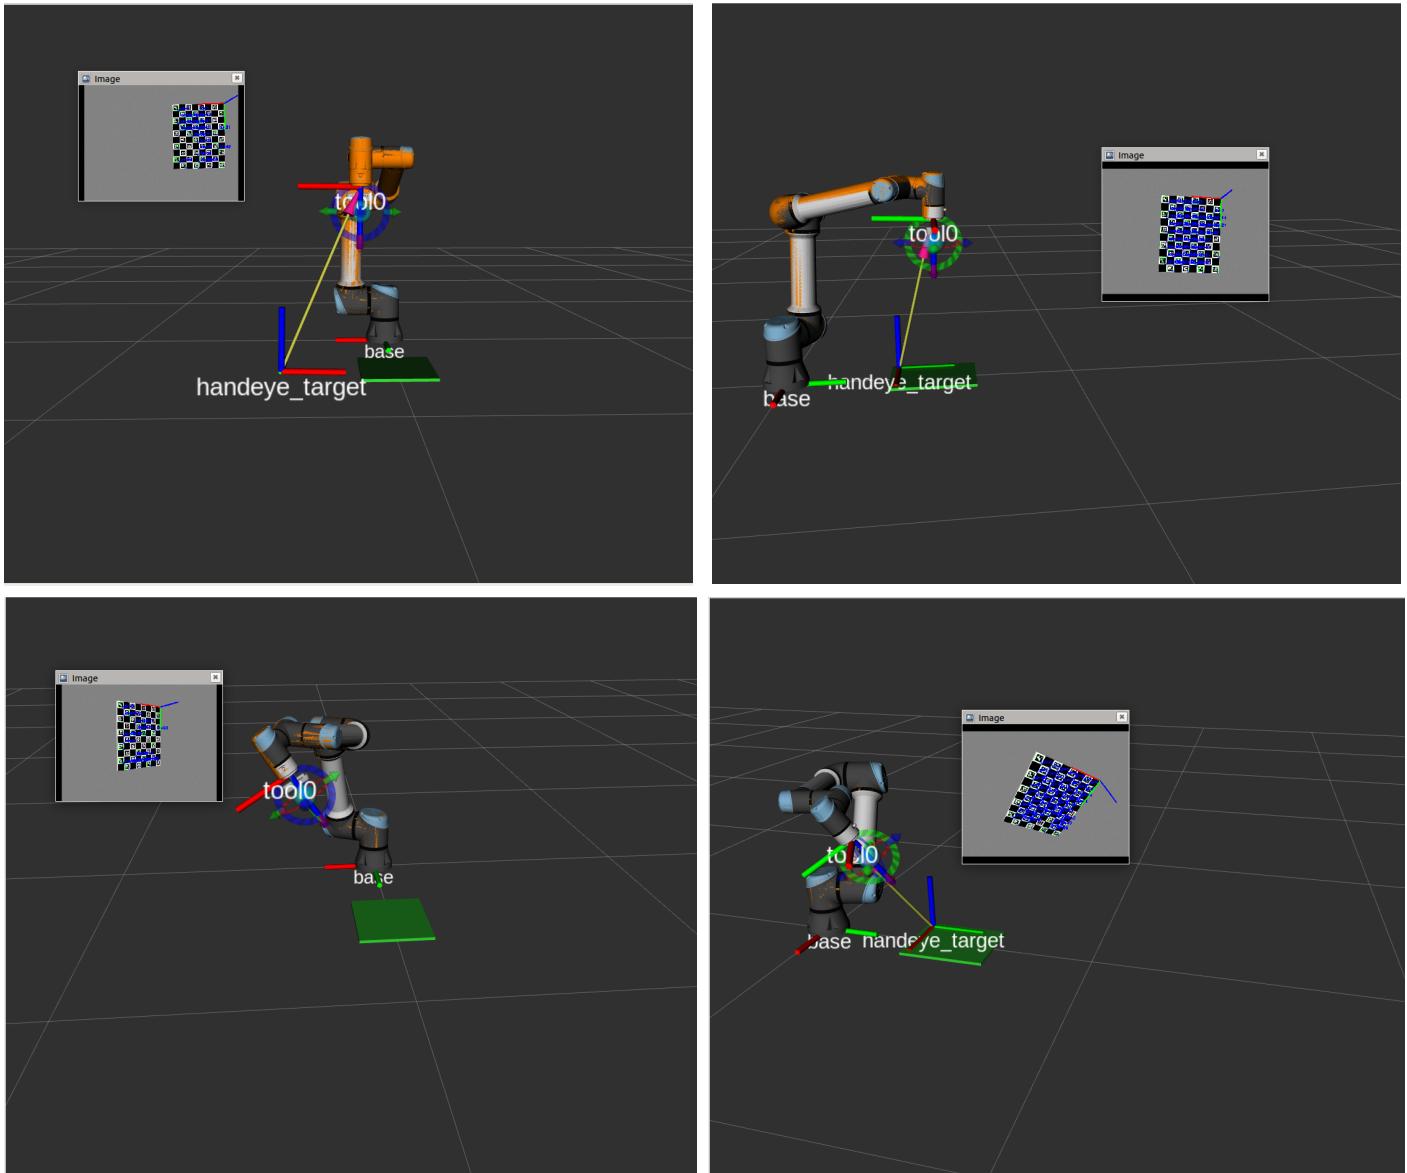


Figure 3 - From the Different to Take Location

Last one is select solve and whis will give the final result of ${}^H T_E$ and three methods result is shown below:

crigroup/Daniilidis 1999

-1	-0.000267438	-0.000332869	0.0172432
0.000267732		-1 -0.000883878	0.0126004
-0.000332632	-0.000883967		0.0300778
0	0	0	1

crigroup/ParkBryan 1994

-1	-0.000267438	-0.000332869	0.0172432
0.000267732		-1 -0.000883878	0.0126004
-0.000332632	-0.000883967		0.0300778
0	0	0	1

crigroup/TsaiLenz 1989

-0.997851	-0.0652883	0.00548093	0.0163231
0.0653405	-0.997814	0.00993659	0.00909533
0.0048202	0.0102734	0.999936	0.035291
0	0	0	1

And the final location coordinate of target and camera are correct as the figure shown below:

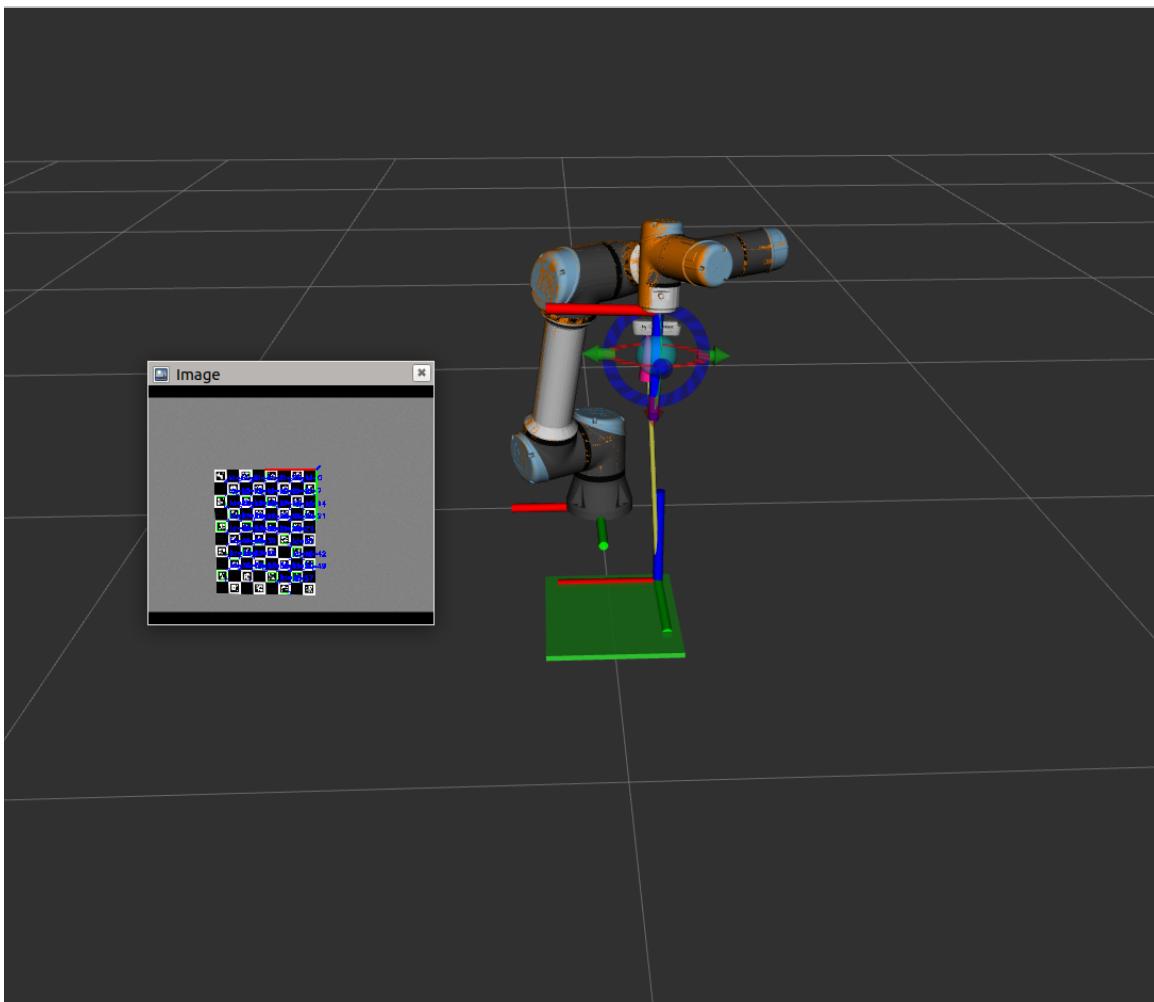


Figure 4 - The Final Coordinate of Target and Camera

Question 2 - Point Cloud Registration

Iterative Closest Point (ICP)

a) Point Cloud Poses

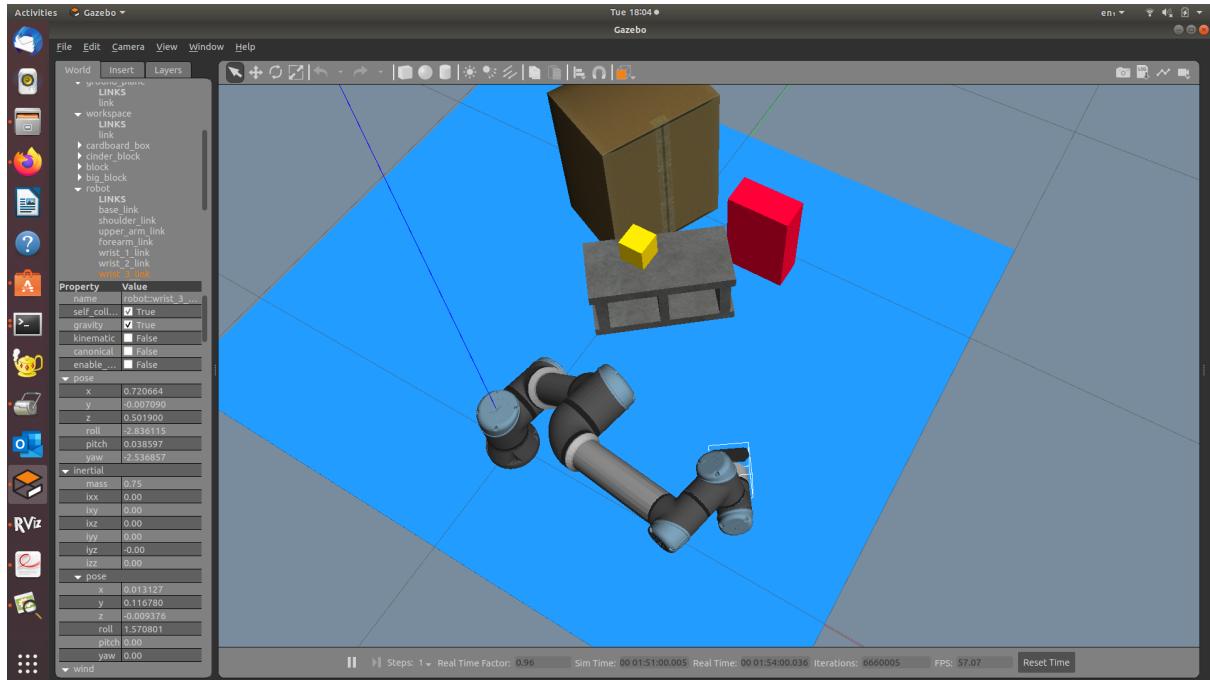


Figure 5 - Pose 1 (End Effector Data)

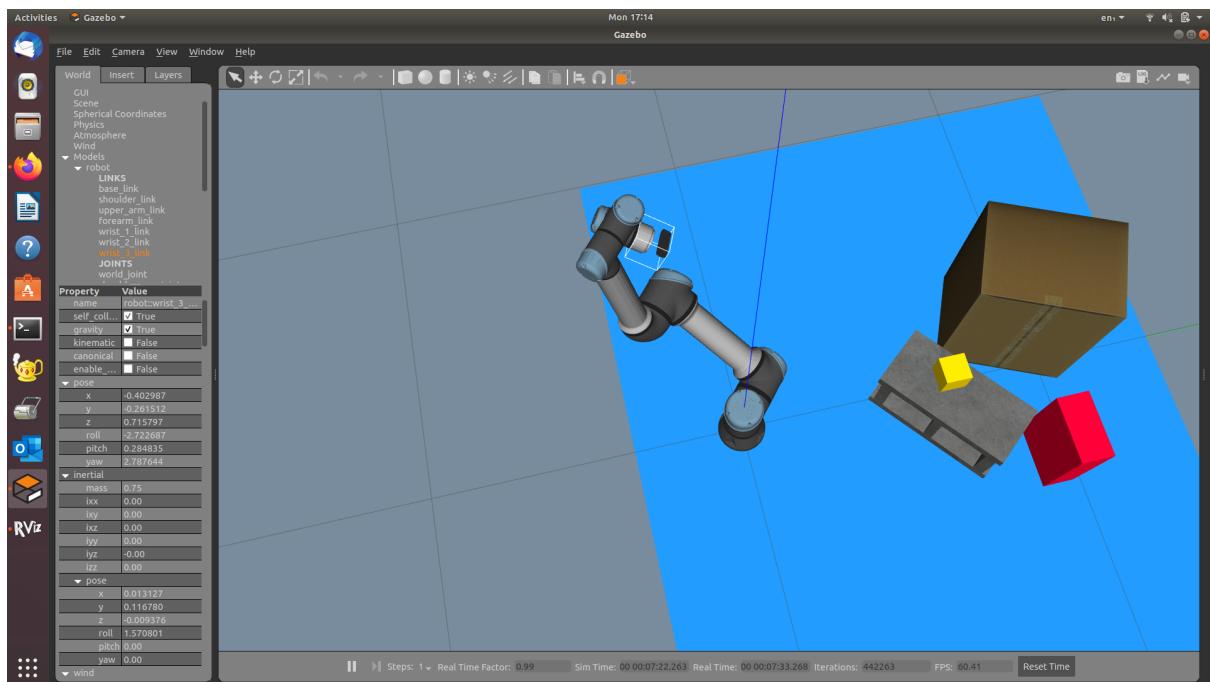


Figure 6 - Pose 2 (End Effector Data)

b) MATLAB Two Point Cloud Visualisation

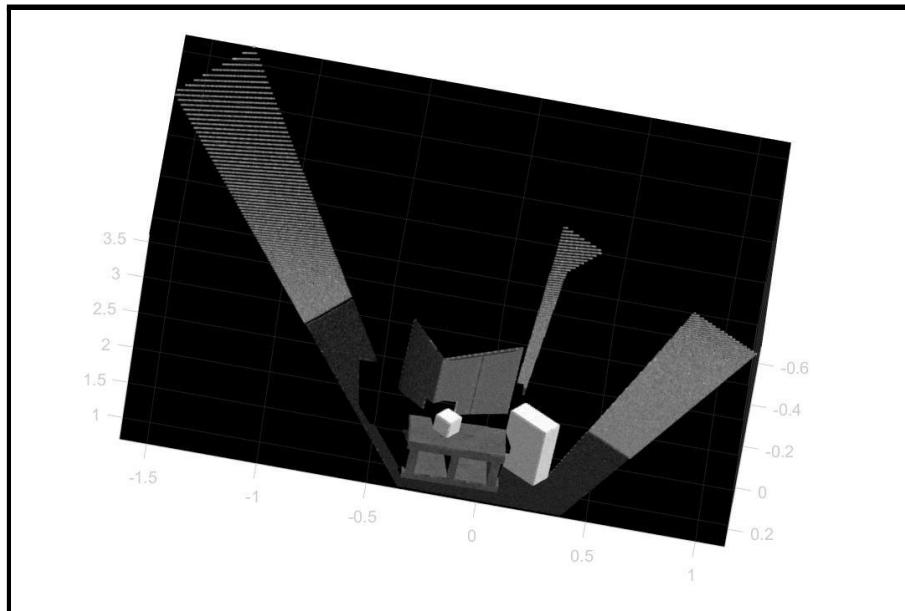


Figure 7 - End-Effector Pose 1 ($T_E^{B^1}$)

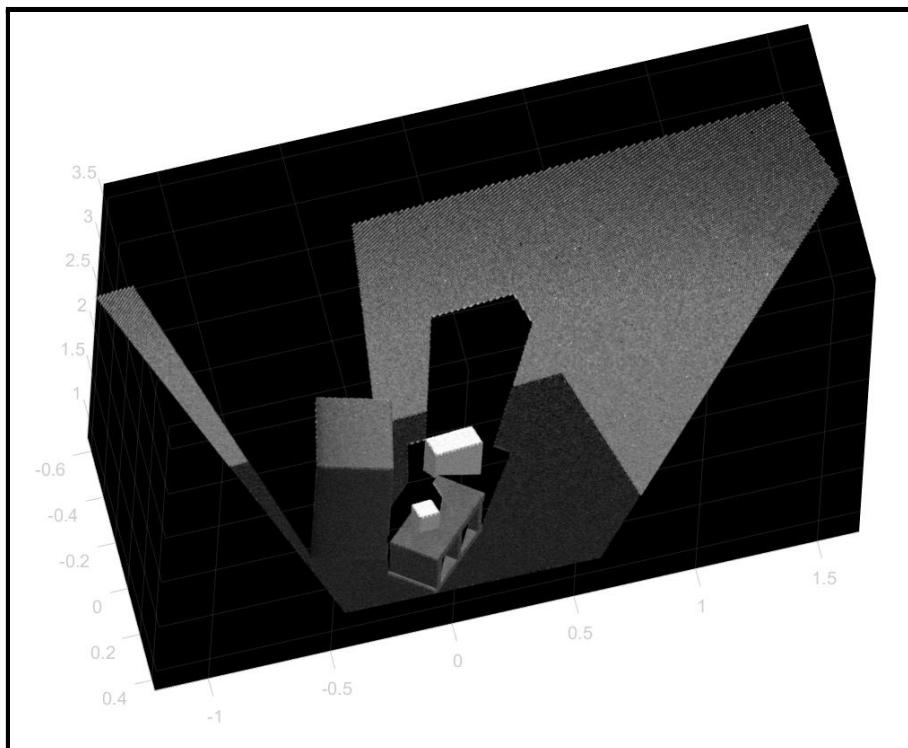


Figure 8 - End-Effector Pose 2 ($T_E^{B^2}$)

c) Results

Motion of the Camera

From Question 1 - it has been assumed that the end-effector to camera transformation

obtained is: $T_E^C =$

$$\begin{matrix} -0.99785 & -0.06529 & 0.005481 & 0.016323 \\ 0.065341 & -0.99781 & 0.009937 & 0.009095 \\ 0.00482 & 0.010273 & 0.999936 & 0.035291 \\ 0 & 0 & 0 & 1 \end{matrix}$$

Furthermore, the poses of the end-effector have a defined transformation matrix as follows (respectively **pose one** and **pose two**):

Pose One $T_E^{B^1} =$

$$\begin{matrix} 0.999 & 0.0443 & 0.0007 & 0.7207 \\ -0.0442 & 0.9978 & 0.0495 & -0.0071 \\ 0.0015 & -0.0495 & 0.9988 & 0.5019 \\ 0 & 0 & 0 & 1 \end{matrix}$$

Pose Two $T_E^{B^2} =$

$$\begin{matrix} 0.9988 & -0.04863 & 0.00497 & -0.40299 \\ 0.048341 & 0.997733 & 0.047499 & -0.26151 \\ -0.00265 & 0.047685 & 0.99888 & 0.715797 \\ 0 & 0 & 0 & 1 \end{matrix}$$

The motion of the camera therefore, can be computed using the following equation:

$$M_2^1 = (T_E^{B^1} \cdot T_C^E)^{-1} \cdot (T_E^{B^2} \cdot T_C^E)$$

Equation 1 - Computation of the Motion of the Camera

This resulting motion of the camera hence is as follows:

0.995692	-0.09286	-0.00349	1.088374
0.092901	0.99005	0.001178	0.400817
0.004968	-0.09627	1.000955	0.191089
0	0	0	1

Point Cloud Registration

The transformation matrix defined utilizing Two Point Cloud Registration within MATLAB (with similar formatting to the above matrices) resulted to be:

0.9727	-0.1886	-0.00639	0.9999
0.09616	0.9729	0.00137	0.2349
0.00843	-0.135	0.9777	0.1609
0	0	0	1

Computation of Errors

Assuming the motion of the camera that was derived from *Equation 1*, is the true expressional representation, the following formula was employed to calculate the error between the true and point cloud registered or ICP derived matrix values:

$$\text{Error} = \frac{\text{ICP Value} - \text{True Matrix Value (Motion of Camera)}}{\text{True Matrix Value (Motion of Camera)}} \times 100$$

Furthermore, the following matrix displays the error between M_2^1 and ICP method, accomplished by subtracting the values of the matrix respectively:

- The **angle of rotation error** can be displayed as:

Rotation	X	Y	Z
X	0.022992	0.095737	0.002901
Y	-0.00326	0.01715	-0.00019
Z	-0.00346	0.038732	0.023255

- The **norm of the translation error** can be displayed as:

Translation

0.0884742
 0.1659169
 0.0301891
translation error = 17.37%

Question 3 - Image Classification using Deep Learning

a) Initial Experimenting

- i) Number of samples in the German Traffic Signs Dataset: 3000
- ii) 1500 of the samples are organized into the training set, while 420 samples are loaded to the validation set and 1080 samples are loaded to the test set.
 - a) Training Data: Data which the network can familiarise itself with the information.
 - b) Test Data: Data which is used by the network for testing.
 - c) Validation Data: Data which is used upon each iteration of the training dataset to investigate the network's performance.
- iii) As shown below, there are **five** classes in the dataset, including: *Stop*, *Right*, *Left*, *Forward* and *Roundabout*. In the three datasets, there are in total 643 *Stop* signs, 562 *Right* signs, 485 *Left* signs, 842 *Forward* signs, and 468 *Roundabout* signs.



Figure 9 - Output Traffic Signs

- iv) Dimensions of the images in pixels: 32×32 pixels

b) Deep Neral Network

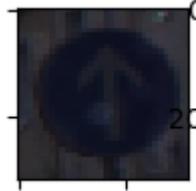
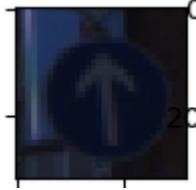
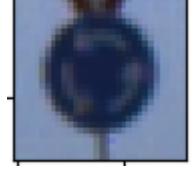
We defined our neural network based on the example provided on the Pytorch Tutorials [1]. During the training loop, we trained the model based on both the training dataset and the validation dataset, with a dataset batch size of 5 and a minibatch of 200, using 5 epochs. The overall performance of the model when experimented on the test dataset reaches an accuracy of 97%.

c) Confusion Matrix

The results of the neural network prediction based on the training dataset is shown below. The model was operated using a batch size of five and a mini-batch of 50. As we can see, the prediction is somewhat accurate apart from several prediction failures. Failed predictions often occur when the image is dark or blurry. The network

cannot predict such images likely because they have a low colour contrast, which is one of the key characteristics during image classification, and hence the network fails to accurately identify their labels.

Table 1 - Confusion Matrix

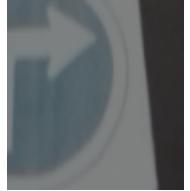
Actual Label	Predicted Label	Image <i>(if it doesn't match)</i>
Stop	Stop	
Forward	Right	
Forward	Forward	
Forward	Forward	
Stop	Stop	
Stop	Stop	
Forward	Right	
Roundabout	Roundabout	
Right	Right	
Left	Left	
Forward	Forward	
Forward	Roundabout	
Roundabout	Left	

d) External Input

We extracted some traffic signs from the rosbag file given in Question 4, preprocessed the images and inputted to the trained model, and below are some of the outputs. There are two failed predictions, as we can see from the table below, the first row and the last row. We observed that all samples in the training datasets had a medium to low brightness, whereas in this failed case the image has a high brightness compared to most of the samples in the training dataset, and hence the model might not be trained to handle bright pictures. This is the same for the case at the last row, where the image has a high brightness but is also captured at an angle, hence more difficult for the model to predict.

Table 2 -

Actual Label	Predicted Label	Image
Right	Roundabout	
Stop	Stop	
Right	Right	
Roundabout	Roundabout	
Right	Right	

Right	Roundabout	
-------	------------	---

Question 4 - Perception for Robots

Navigate Duckietown

a) Process of Navigation

The process of identifying the traffic signs through the manufactured model town (Duckietown) was achieved by firstly registering and detecting a **red** line. This line acts as an indicator within the system, highlighting to the program that an oncoming traffic sign on the surface will be present and requires detection. Furthermore, the red line coordinates indicate to the system to trigger the binary **white** colour detection. When this colour is detected, its coordinates can be found using further OpenCV functions and the frame of that traffic sign can be grabbed. This is accomplished alongside the red line disappearing from the frames, triggering the frame capture of the traffic sign to be centralised and fully present in the saving of that particular image/frame.

Finally, this frame must undergo some further processing prior to it being applied to the neural network - this includes transforming the trapezoidal like traffic sign (since the camera is slightly above the surface and is in a parallel perspective to the surface) to a more square image and reducing the captured image to 32×32 a pixel image.

From this stage the neural network can then detect the label, predict the traffic sign and display the image, label and traffic sign as an output.

Upon testing the program, many alterations were made to improve the performance of the system in terms of efficiency and accuracy - this included reducing the contrast and brightness of the images prior to detection of the red line. It was evident that this reduced the glare present on the red line and hence, the number of frames that are saved will be kept to a minimal (as there will be a more clear indication as to when the line is present verses when it is not - contrast between the red, black and white aspects of the images).

Alike, to reduce the number of frames captured to a minimal or ideally only one frame for each traffic sign, a timer function was planned to be employed but was not due to time restrictions. This would limit the amount of captures for each traffic sign and hence, increase the accuracy of the system especially with Question 5.

Question 5 - Demo

- Please refer to video attachment entitled:
 - *Question 5*

Appendix

Codes for Question 3-4

```
1  #!/usr/bin/env python
2  # -*- coding: utf-8 -*-
3  import pickle
4  import matplotlib.pyplot as plt
5  import cv2
6  import torch
7  import time
8  import PIL
9  from PIL import Image
10 import numpy as np
11 import tensorflow as tf
12 import torch.nn as nn
13 import torch.nn.functional as functional
14
15 from torch import nn, optim
16 from torch.autograd import Variable
17 from torch.utils.data import Dataset, DataLoader
18 import torchvision
19 from torchvision import transforms
20 from torchvision.transforms import ToTensor
21 from torchvision.transforms import functional as fn
22
23
24 import rosbag
25 import roslib    #roslib.load_manifest(PKG)
26 import rospy
27
28 from cv_bridge import CvBridge
29 from sensor_msgs.msg import Image
30 from scipy.ndimage import filters
31
32 font = cv2.FONT_HERSHEY_SIMPLEX
33
34 training_file = "/home/ob1-knb/ur5/src/traindata.pkl"
35 validation_file = "/home/ob1-knb/ur5/src/validdata.pkl"
36 testing_file = "/home/ob1-knb/ur5/src/testdata.pkl"
37 xtra_file = "/home/ob1-knb/ur5/src/extradata.pkl"
38
39 with open(training_file, mode='rb') as f:
40     train = pickle.load(f)
41 with open(validation_file, mode='rb') as f:
42     valid = pickle.load(f)
43 with open(testing_file, mode='rb') as f:
44     test = pickle.load(f)
```

```
45 with open(xtra_file, mode='rb') as f:
46     xtra = pickle.load(f)
47
48
49 x_train, y_train = train['features'], train['labels']
50 x_valid, y_valid = valid['features'], valid['labels']
51 x_test, y_test = test['features'], test['labels']
52 x_xtra, y_xtra = xtra['features'], xtra['labels']
53
54 # Custom Dataloader Class
55 class TrafficSignsDataset(Dataset):
56     def __init__(self, features, labels, transform=None):
57         self.features = features
58         self.labels = labels
59         self.transform = transform
60
61     def __len__(self):
62         return len(self.features)
63
64     def __getitem__(self, idx):
65         if torch.is_tensor(idx):
66             idx = idx.tolist()
67
68         x = self.features[idx]
69         y = self.labels[idx]
70
71         if self.transform:
72             x = self.transform(x)
73             y = y
74
75         return x, y
76
77
78 # Define custom transforms to apply to the images
79 # -----
80 # Here we are applying a normalization to bring all pixel values between 0 and 1
81 # and converting the data to a pytorch tensor
82 custom_transform = transforms.Compose([transforms.ToTensor(),
83                                     transforms.Normalize((0.5,0.5,0.5), (0.5,0.5,0.5))])
84
85 # Define the dataloaders
86 # -----
87 # Dataloaders help in managing large amounts of data efficiently
```

```

88 batch_size = 5
89
90 train_dataset = TrafficSignsDataset(x_train, y_train,
91                                     custom_transform)
92 train_loader = DataLoader(train_dataset, batch_size=batch_size, shuffle=True)
93
94 validation_dataset = TrafficSignsDataset(x_valid, y_valid,
95                                         custom_transform)
96 validation_loader = DataLoader(validation_dataset, shuffle=True)
97
98 test_dataset = TrafficSignsDataset(x_test, y_test,
99                                     custom_transform)
100 test_loader = DataLoader(test_dataset, batch_size=batch_size, shuffle=True)
101
102 xtra_dataset = TrafficSignsDataset(x_xtra, y_xtra,
103                                     custom_transform)
104 xtra_loader = DataLoader(xtra_dataset, batch_size=batch_size, shuffle=True)
105
106
107
108 classes = ('Stop', 'Right', 'Left', 'Forward', 'Roundabout')
109
110 def imshow(img):
111     img = img / 2 + 0.5
112
113 dataiter = iter(train_loader)
114 images, labels = dataiter.next()
115
116 print("Number of training samples is {}".format(len(train_dataset)))
117 print("Number of test samples is {}".format(len(test_dataset)))
118 print("Number of validation samples is {}".format(len(validation_dataset)))
119
120 print("Batch size is {}".format(len(images)))
121 print("Size of each image is {}".format(images[0].shape))
122
123 print("The labels in this batch are: {}".format(labels))
124 print("These correspond to the classes: {}, {}, {}, {}, {}, {}".
125       format(classes[labels[0]], classes[labels[1]],
126             classes[labels[2]], classes[labels[3]], classes[labels[4]]))
127
128 #####
129 # Step 2: Building your Neural Network
130 #
131

```

```
132 class Network(nn.Module):
133     def __init__(self):
134         self.output_size = 5    # 5 classes
135
136         super(Network, self).__init__()
137         self.conv1 = nn.Conv2d(3, 6, 5)           # 2D convolution
138         self.pool = nn.MaxPool2d(2, 2)          # max pooling
139         self.conv2 = nn.Conv2d(6, 16, 5)          # 2D convolution
140         self.fc1 = nn.Linear(16 * 5 * 5, 120)    # Fully connected layer
141         self.fc2 = nn.Linear(120, 84)           # Fully connected layer
142         self.fc3 = nn.Linear(84, self.output_size) # Fully connected layer
143
144     def forward(self, x):
145         """Define the forward pass."""
146         x = self.pool(functional.relu(self.conv1(x)))
147         x = self.pool(functional.relu(self.conv2(x)))
148         x = x.view(-1, 16 * 5 * 5)
149         x = functional.relu(self.fc1(x))
150         x = functional.relu(self.fc2(x))
151         x = self.fc3(x)
152         return x
153
154 net = Network()
155
156 ##### Step 3: Defining loss function and optimizer #####
157 # Step 3: Defining loss function and optimizer
158 #-----
159
160 criterion = nn.CrossEntropyLoss()
161 optimizer = optim.SGD(net.parameters(), lr=0.001, momentum=0.9)
162
163 ##### Step 4: Main training loop #####
164 # Step 4: Main training loop
165 #-----
166
167 for epoch in range(5):
168     running_loss = 0.0
169     val_loss = 0.0
170
171     for i, data in enumerate(validation_loader, 0):
172
173         inputs, labels = data
174         optimizer.zero_grad()
```

```

175     predicted_labels = net(inputs)
176     loss = criterion(predicted_labels, labels)
177     loss.backward()
178     optimizer.step()
179
180     val_loss += loss.item()
181     if i % 200 == 199:    # print every 200 mini-batches
182         print('Epoch: %d, %5d val loss: %.3f' % (epoch + 1, i + 1, val_loss / 50))
183         val_loss = 0.0
184
185     for i, data in enumerate(train_loader, 0):
186
187         inputs, labels = data
188         optimizer.zero_grad()
189         predicted_labels = net(inputs)
190         loss = criterion(predicted_labels, labels)
191         loss.backward()
192         optimizer.step()
193
194         running_loss += loss.item()
195         if i % 200 == 199:
196             print('Epoch: %d, %5d loss: %.3f' % (epoch + 1, i + 1, running_loss / 50))
197             running_loss = 0.0
198
199     print('Finished Training.')
200     torch.save(net.state_dict(), "traffic_signs_classifier.pt")
201     print('Saved model parameters to disk.')
202
203 #####
204 # Step 5: Duckietown Experiment
205 #-----
206
207 model = Network()
208 model.load_state_dict(torch.load('traffic_signs_classifier.pt'))
209 model.eval()
210
211 bag_file = '/home/ob1-knb/ur5/src/bag1.bag'
212 bag = rosbag.Bag(bag_file, "r")
213
214 bridge = CvBridge()
215 bag_data = bag.read_messages('/raspicam_node/image/compressed')
216
217 for topic, msg, t in bag_data:
218     cv_image = bridge.compressed_imgmsg_to_cv2(msg, "bgr8")

```

```

217 for topic, msg, t in bag_data:
218     cv_image = bridge.compressed_imgmsg_to_cv2(msg, "bgr8")
219
220     dst = cv2.rotate(cv_image, cv2.ROTATE_180)
221     original = dst
222     alpha = 0.5
223     beta = 0
224     res = cv2.convertScaleAbs(dst, alpha=alpha, beta=beta)
225     HSVImage_imp = cv2.cvtColor(res, cv2.COLOR_BGR2HSV)
226     Low_red_range = np.array([140, 50, 50])
227     High_red_range = np.array([180, 255, 255])
228     Low_white_range = np.array([100, 43, 46])
229     High_white_range = np.array([124, 255, 255])
230
231     mask_red = cv2.inRange(HSVImage_imp, Low_red_range, High_red_range)
232     mask_red = cv2.medianBlur(mask_red, 7)
233
234     mask_red, contours2, hierarchy2 = cv2.findContours(mask_red, cv2.RETR_EXTERNAL, cv2.CHAIN_APPROX_NONE)
235
236     for cnt2 in contours2:
237         (x2, y2, w2, h2) = cv2.boundingRect(cnt2)
238
239         if (y2>308 and y2<315) and w2 > 3*h2 :
240
241             crop_img = HSVImage_imp[220:300,130:490]
242             mask_white = cv2.inRange(crop_img,Low_white_range,High_white_range)
243             mask_white = cv2.medianBlur(mask_white, 7) # 中值滤波
244             mask_red1 = cv2.inRange(crop_img,Low_red_range,High_red_range)
245             mask_red1 = cv2.medianBlur(mask_red1, 7) # 中值滤波
246             mask = cv2.bitwise_or(mask_red1,mask_white)
247             mask_white, contours1, hierarchy1 = cv2.findContours(mask, cv2.RETR_EXTERNAL, cv2.CHAIN_APPROX_NONE)
248
249             for cnt1 in contours1:
250                 (x, y, w, h) = cv2.boundingRect(cnt1)
251                 crop_img_1 = crop_img[y:y+60,x:x+120]
252                 RGB_img_1 = cv2.cvtColor(crop_img_1, cv2.COLOR_HSV2BGR)
253
254                 resize_img = cv2.resize(RGB_img_1,(32,32))
255
256
257                 resize_img = resize_img[:, :, [2,1,0]]
258                 # resize_img = resize_img[:, :, ::-1]

```

```

260     RGB_img_1 = cv2.resize(RGB_img_1,(100,100))
261
262     cv2.imshow("Image window",RGB_img_1)
263
264     x = ToTensor()(resize_img)
265     x = transforms.Normalize((0.5,0.5,0.5 ), (0.5,0.5,0.5))(x)
266     x.unsqueeze_(0)
267
268     outputs = model(x)
269     sm = nn.Softmax(dim=1)
270     sm_outputs = sm(outputs)
271
272     probs, index = torch.max(sm_outputs, dim=1)
273     for p, i in zip(probs, index):
274         print('{}'.format(classes[i]))
275         sign = classes[i]
276         cv2.putText(original, '{}'.format(sign), (240,220), font, 2, (0,0,0))
277         time.sleep(0.5)
278
279     cv2.imshow("Image window", original)
280
281
282     cv2.waitKey(3)
283
284 #####
285 # Step 6: Evaluate the accuracy of your network on the test dataset
286 #-----
287
288 dataiter = iter(train_loader)
289 images, labels = dataiter.next()
290
291 fig, ax = plt.subplots(1, len(images))
292 for id, image in enumerate(images):
293     # convert tensor back to numpy array for visualization
294     ax[id].imshow((image / 2 + 0.5).numpy().transpose(1,2,0))
295     ax[id].set_title(classes[labels[id]])
296 plt.show()
297 print(' '.join('%5s' % classes[labels[j]] for j in range(batch_size)))
298
299 outputs = model(images)
300

```

```

301 sm = nn.Softmax(dim=1)
302 sm_outputs = sm(outputs)
303
304 probs, index = torch.max(sm_outputs, dim=1)
305 for p, i in zip(probs, index):
306     print('True label {0}, Predicted label {0} - {1:.4f}'.format(classes[i], p))
307
308 correct = 0
309 total = 0
310 with torch.no_grad():
311     for data in validation_loader:
312         images, labels = data
313         outputs = model(images)
314         _, predicted = torch.max(outputs.data, 1)
315         total += labels.size(0)
316         correct += (predicted == labels).sum().item()
317
318 print('Correct predictions: {}'.format(correct))
319 print('Total predicted: {}'.format(total))
320 print('Number per class: ')
321 print(stop)
322 print(right)
323 print(left)
324 print(forward)
325 print(round)
326 print('Accuracy of the network on the 10000 test images: %d %%' % (
327     100 * correct / total))
328

```

Codes for Question 5

```

#!/usr/bin/env python3
#!/usr/bin/env python2
# -*- coding: utf-8 -*-
import pickle
import matplotlib.pyplot as plt
import cv2
import torch
import time
from PIL import Image
import numpy as np
import tensorflow as tf
import torch.nn as nn
import torch.nn.functional as functional
import sys, time
from torch import nn, optim
from torch.autograd import Variable
from torch.utils.data import Dataset, DataLoader
import torchvision
from torchvision import transforms

```

```
from torchvision.transforms import ToTensor
from scipy.ndimage import filters
import rosbag
import roslib #roslib.load_manifest(PKG)
import rospy
from sensor_msgs.msg import CompressedImage
from cv_bridge import CvBridge
from sensor_msgs.msg import Image
from scipy.ndimage import filters
from geometry_msgs.msg import Twist
PI = 3.1415926535897
```

```
training_file = "/home/max/assignment2/src/traindata.pkl"
validation_file = "/home/max/assignment2/src/validdata.pkl"
testing_file = "/home/max/assignment2/src/testdata.pkl"
```

```
with open(training_file, mode='rb') as f:
    train = pickle.load(f)
with open(validation_file, mode='rb') as f:
    valid = pickle.load(f)
with open(testing_file, mode='rb') as f:
    test = pickle.load(f)
```

```
x_train, y_train = train['features'], train['labels']
x_valid, y_valid = valid['features'], valid['labels']
x_test, y_test = test['features'], test['labels']
```

```
# Custom Dataloader Class
```

```
class TrafficSignsDataset(Dataset):
    def __init__(self, features, labels, transform=None):
        self.features = features
        self.labels = labels
        self.transform = transform

    def __len__(self):
        return len(self.features)

    def __getitem__(self, idx):
        if torch.is_tensor(idx):
            idx = idx.tolist()

        x = self.features[idx]
        y = self.labels[idx]

        if self.transform:
            x = self.transform(x)
```

```

y = y

return x, y

# Define custom transforms to apply to the images
# -----
# Here we are applying a normalization to bring all pixel values between 0 and 1
# and converting the data to a pytorch tensor
custom_transform = transforms.Compose([transforms.ToTensor(),
                                      transforms.Normalize((0.5,0.5,0.5 ), (0.5,0.5,0.5))])

# Define the dataloaders
# -----
# Dataloaders help in managing large amounts of data efficiently
batch_size = 5

train_dataset = TrafficSignsDataset(x_train, y_train,
                                    custom_transform)
train_loader = DataLoader(train_dataset, batch_size=batch_size, shuffle=True)

validation_dataset = TrafficSignsDataset(x_valid, y_valid,
                                         custom_transform)
validation_loader = DataLoader(validation_dataset, shuffle=True)

test_dataset = TrafficSignsDataset(x_test, y_test,
                                   custom_transform)
test_loader = DataLoader(test_dataset, batch_size=batch_size, shuffle=True)

classes = ('Stop', 'Right', 'Left', 'Forward', 'Roundabout')

class Network(nn.Module):
    def __init__(self):
        super().__init__()
        self.output_size = 5 # 5 classes

        self.conv1 = nn.Conv2d(3, 6, 5) # 2D convolution
        self.pool = nn.MaxPool2d(2, 2) # max pooling
        self.conv2 = nn.Conv2d(6, 16, 5) # 2D convolution
        self.fc1 = nn.Linear(16 * 5 * 5, 120) # Fully connected layer
        self.fc2 = nn.Linear(120, 84) # Fully connected layer
        self.fc3 = nn.Linear(84, self.output_size) # Fully connected layer

    def forward(self, x):
        """Define the forward pass."""

```

```

x = self.pool(functional.relu(self.conv1(x)))
x = self.pool(functional.relu(self.conv2(x)))
x = x.view(-1, 16 * 5 * 5)
x = functional.relu(self.fc1(x))
x = functional.relu(self.fc2(x))
x = self.fc3(x)
return x

net = Network()

#####
#####

#####
#####
# Step 3: Defining loss function and optimizer
#-----

criterion = nn.CrossEntropyLoss()
optimizer = optim.SGD(net.parameters(), lr=0.001, momentum=0.9)

#####
#####

# Step 4: Main training loop
#-----


for epoch in range(6): # we are using 5 epochs. Typically 100-200
    running_loss = 0.0
    for i, data in enumerate(train_loader, 0):
        # get the inputs
        inputs, labels = data

        # zero the parameter gradients
        optimizer.zero_grad()

        # Perform forward pass and predict labels
        predicted_labels = net(inputs)

        # Calculate loss
        loss = criterion(predicted_labels, labels)

        # Perform back propagation and compute gradients
        loss.backward()

        # Take a step and update the parameters of the network
        optimizer.step()

        # print statistics

```

```

running_loss += loss.item()
if i % 300 == 299:      # print every 200 mini-batches
    print('Epoch: %d, %5d loss: %.3f' % (epoch + 1, i + 1, running_loss / 50))
running_loss = 0.0

print('Finished Training.')
torch.save(net.state_dict(), "traffic_signs_classifier.pt")
torch.save(net, '/home/max/assignment2/cifar_2.pt')
print('Saved model parameters to disk.')

#####
#####
# Step 5: Evaluate the accuracy of your network on the test dataset
#-----

model = Network()
model.load_state_dict(torch.load('traffic_signs_classifier.pt'))
model.eval()

class image_feature:

    def __init__(self):
        # subscribed Topic
        self.subscriber = rospy.Subscriber('/raspicam_node/image/compressed',
                                         CompressedImage, self.callback, queue_size=1)

    def callback(self, ros_data):
        velocity_publisher = rospy.Publisher('/turtle1/cmd_vel', Twist, queue_size=1)
        angular_speed = 90*2*PI/360
        relative_angle = 90*2*PI/360
        vel_msg = Twist()
        vel_msg.linear.x = 0.1
        vel_msg.linear.y = 0
        vel_msg.linear.z = 0
        vel_msg.angular.x = 0
        vel_msg.angular.y = 0
        vel_msg.angular.z = 0
        velocity_publisher.publish(vel_msg)

        font = cv2.FONT_HERSHEY_SIMPLEX
        ##### direct conversion to CV2 #####
        np_arr = np.fromstring(ros_data.data, np.uint8)
        image_np = cv2.imdecode(np_arr, cv2.IMREAD_COLOR) #original image
        image_np = cv2.rotate(image_np, cv2.ROTATE_180)      #rotate
        alpha = 0.7    # change the contrast
        beta = 0

        # eliminate glare

```

```

R, G, B = cv2.split(image_np)
output1_R = cv2.equalizeHist(R)
output1_G = cv2.equalizeHist(G)
output1_B = cv2.equalizeHist(B)
equ = cv2.merge((output1_R, output1_G, output1_B))
res = np.hstack((image_np, equ)) #stacking images side-by-side

# change the contrast use to test the redline
adjusted = cv2.convertScaleAbs(image_np, alpha=alpha, beta=beta)
HSVImage = cv2.cvtColor(adjusted, cv2.COLOR_BGR2HSV) # for redline detection
HSVImage_imp = cv2.cvtColor(res, cv2.COLOR_BGR2HSV) #for traffic signal
recognition

Low_red_range = np.array([140,50,50])
High_red_range = np.array([180,255,255])
Low_white_range = np.array([80,35,70])
High_white_range = np.array([135,230,255])

mask_red = cv2.inRange(HSVImage, Low_red_range, High_red_range)
mask_red = cv2.medianBlur(mask_red, 7) # Median filer
contours2, hierarchy2 = cv2.findContours(mask_red, cv2.RETR_EXTERNAL,
cv2.CHAIN_APPROX_NONE)

for cnt2 in contours2:
    (x2, y2, w2, h2) = cv2.boundingRect(cnt2)
    if w2>3*h2 and (y2==308 or y2==307 or y2 == 314): # take one flame each signal
        crop_img = HSVImage_imp[200:310,190:490]
        mask_white = cv2.inRange(crop_img, Low_white_range, High_white_range)
        mask_white = cv2.medianBlur(mask_white, 7) # Median filer
        mask_red1 = cv2.inRange(crop_img, Low_red_range, High_red_range)
        mask_red1 = cv2.medianBlur(mask_red1, 7) # Median filer
        mask = cv2.bitwise_or(mask_red1, mask_white)
        cv2.imshow('mask_white', mask)
        contours1, hierarchy1 = cv2.findContours(mask, cv2.RETR_EXTERNAL,
cv2.CHAIN_APPROX_NONE)
        for cnt1 in contours1:
            (x, y, w, h) = cv2.boundingRect(cnt1)
            crop_img_1 = crop_img[y:y+60,x:x+120]
            BGR_img_1 = cv2.cvtColor(crop_img_1, cv2.COLOR_HSV2BGR)
            cv2.imwrite('traffic_image_original_1.png', BGR_img_1)
            resize_img = cv2.resize(BGR_img_1, (32,32))
            cv2.imwrite('traffic_image_2.png', resize_img)

# use for model detect
resize_img_1 = resize_img[:, :, [2, 1, 0]]
x = ToTensor()(resize_img_1)
x = transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))(x)
x.unsqueeze_(0)

```

```

outputs = model(x)
sm = nn.Softmax(dim=1)
sm_outputs = sm(outputs)
_, predicted = torch.max(sm_outputs, dim=1)
print('Predicted: ', ' '.join('%5s' % classes[predicted[j]]
                               for j in range(1)))
sign = classes[predicted]
cv2.putText(image_np,format(sign), (240,220),font,2, (0,0,0),2, cv2.LINE_AA)

if sign == "Stop":
    vel_msg.linear.x = 0
    vel_msg.linear.y = 0
    vel_msg.linear.z = 0
    vel_msg.angular.x = 0
    vel_msg.angular.y = 0
    vel_msg.angular.z = 0
    velocity_publisher.publish(vel_msg)

elif sign == "Right":
    current_angle = 0
    t0 = rospy.Time.now().to_sec()
    while(current_angle < relative_angle):
        vel_msg.angular.z = -abs(angular_speed)
        velocity_publisher.publish(vel_msg)
        t1 = rospy.Time.now().to_sec()
        current_angle = angular_speed*(t1-t0)

elif sign == "Left":
    current_angle = 0
    t0 = rospy.Time.now().to_sec()
    while(current_angle < relative_angle):
        vel_msg.angular.z = abs(angular_speed)
        velocity_publisher.publish(vel_msg)
        t1 = rospy.Time.now().to_sec()
        current_angle = angular_speed*(t1-t0)

elif sign == "Roundabout":
    current_angle = 0
    t0 = rospy.Time.now().to_sec()
    while(current_angle < relative_angle*2):
        vel_msg.angular.z = -abs(angular_speed)
        velocity_publisher.publish(vel_msg)
        t1 = rospy.Time.now().to_sec()
        current_angle = angular_speed*(t1-t0)

cv2.imshow('cv_img', image_np)
cv2.waitKey(1)

```

```
def main(args):
    ic = image_feature()
    rospy.init_node('image_feature', anonymous=True)
    try:
        rospy.spin()
    except KeyboardInterrupt:
        print ("Shutting down ROS Image feature detector module")
        cv2.destroyAllWindows()

if __name__ == '__main__':
    main(sys.argv)
```